Direct-buffer GPU + NPU prefill FFN partition

Fresh all-GPU FFN baseline · GPU + NPU, serial · GPU + NPU, concurrent

p2560 first
M=512, H=4096, I=14336 · 10 samples/mode/partition · BFP16

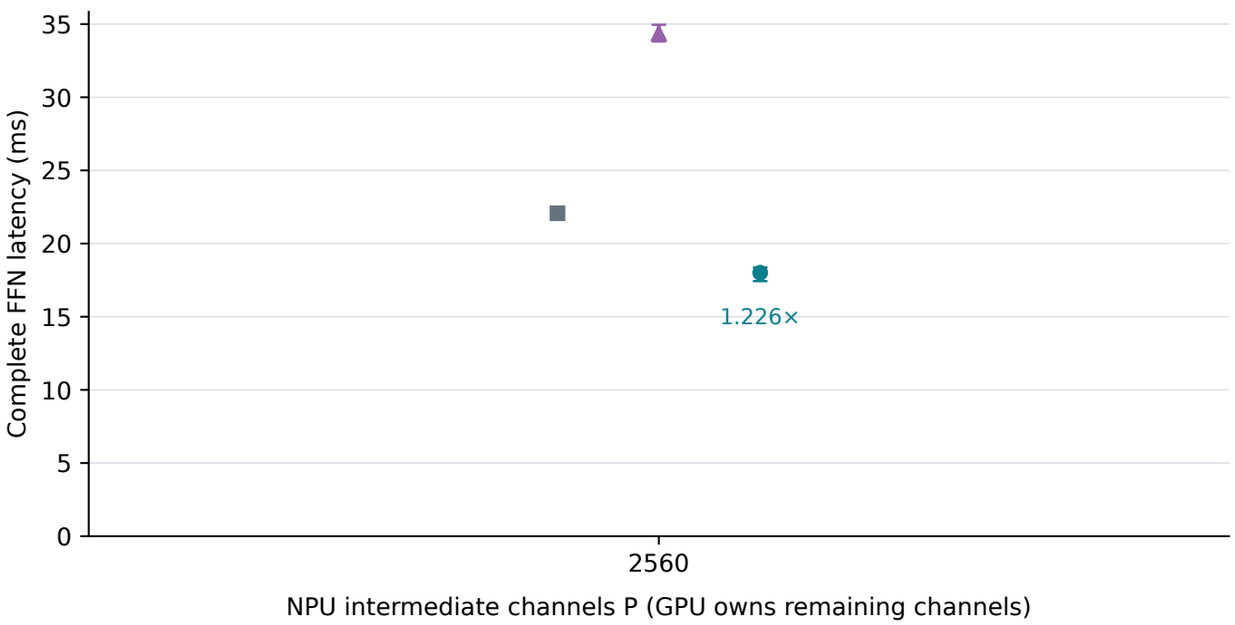

sweep
M=512, H=4096, I=14336 · 20 samples/mode/partition · BFP16

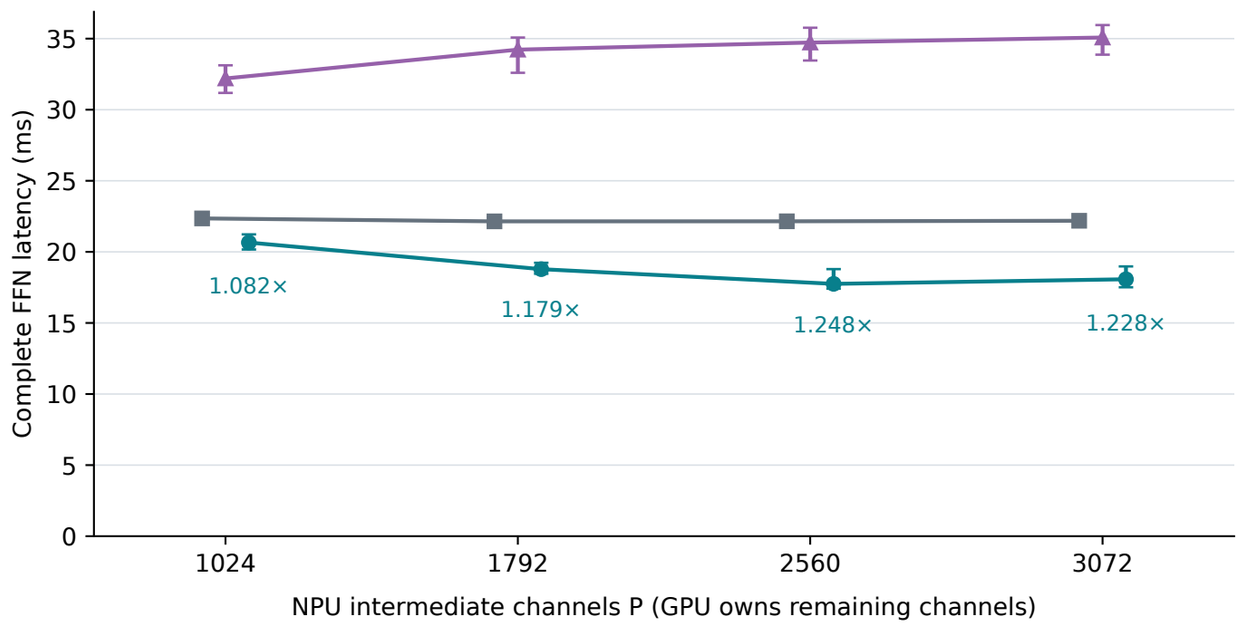

refine
M=512, H=4096, I=14336 · 20 samples/mode/partition · BFP16

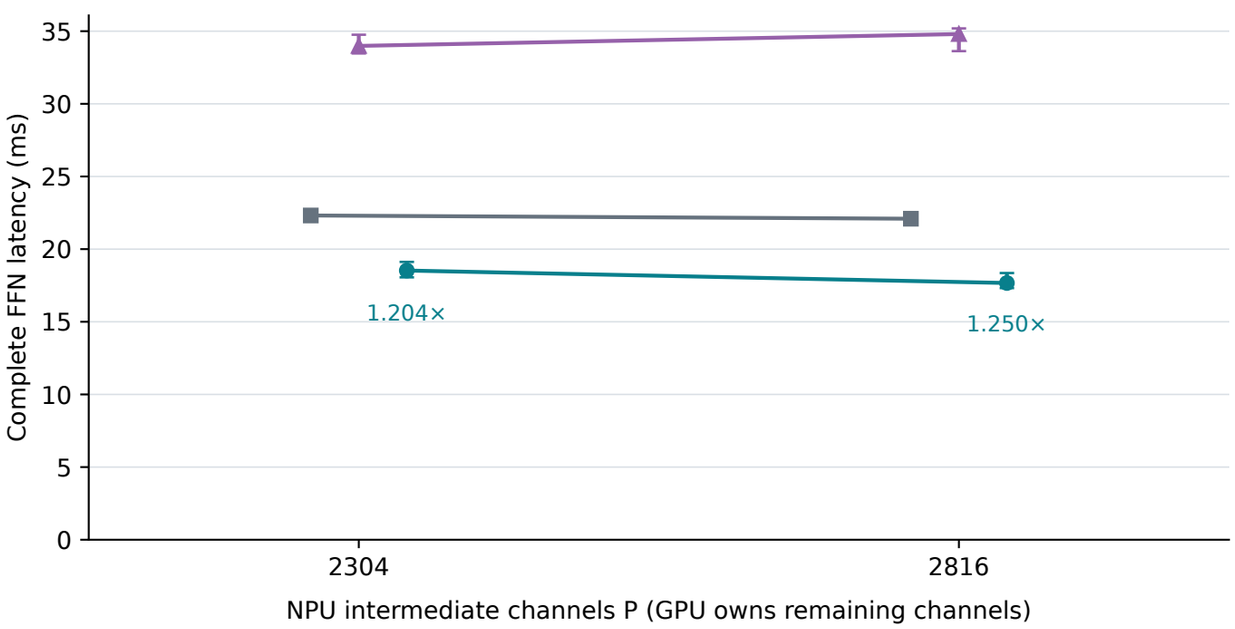

confirm a
M=512, H=4096, I=14336 · 50 samples/mode/partition · BFP16

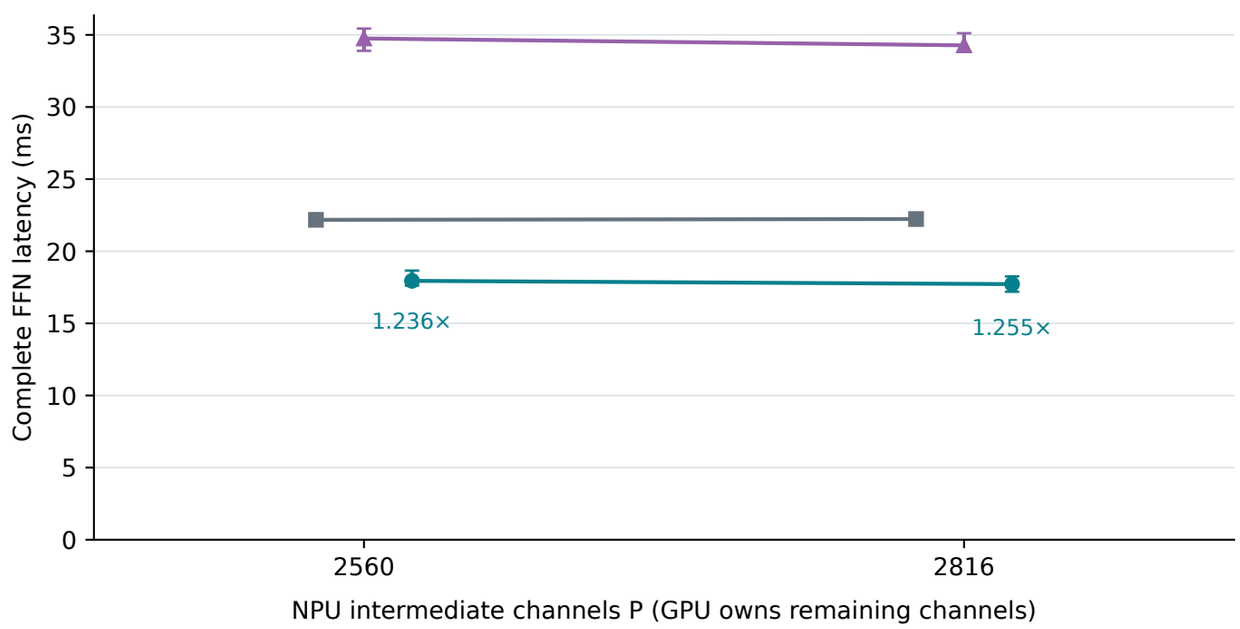

confirm b
M=512, H=4096, I=14336 · 50 samples/mode/partition · BFP16

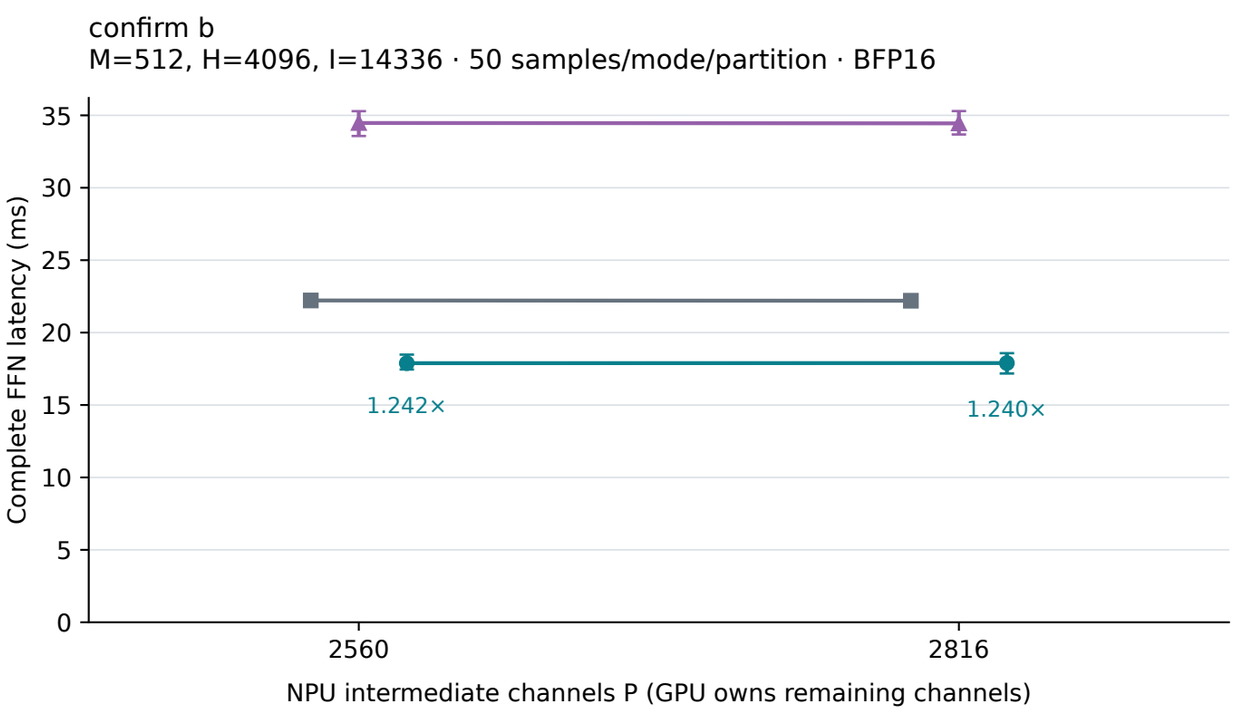

Dots: median; bars: p10–p90. × labels: fresh GPU median / concurrent median. Each panel is a separate run; samples are never pooled. Timing includes host staging, BO I/O, waits and partial-output reduction. Isolated FFN only; no whole-model/TTFT claim. † Original tolerance deviation.