W4A16 BFP16 FFN: zero-point conversion and host readback optimization

M512 / H4096 / I14336 · shared GPU–NPU FFN partition · final sum on GPU

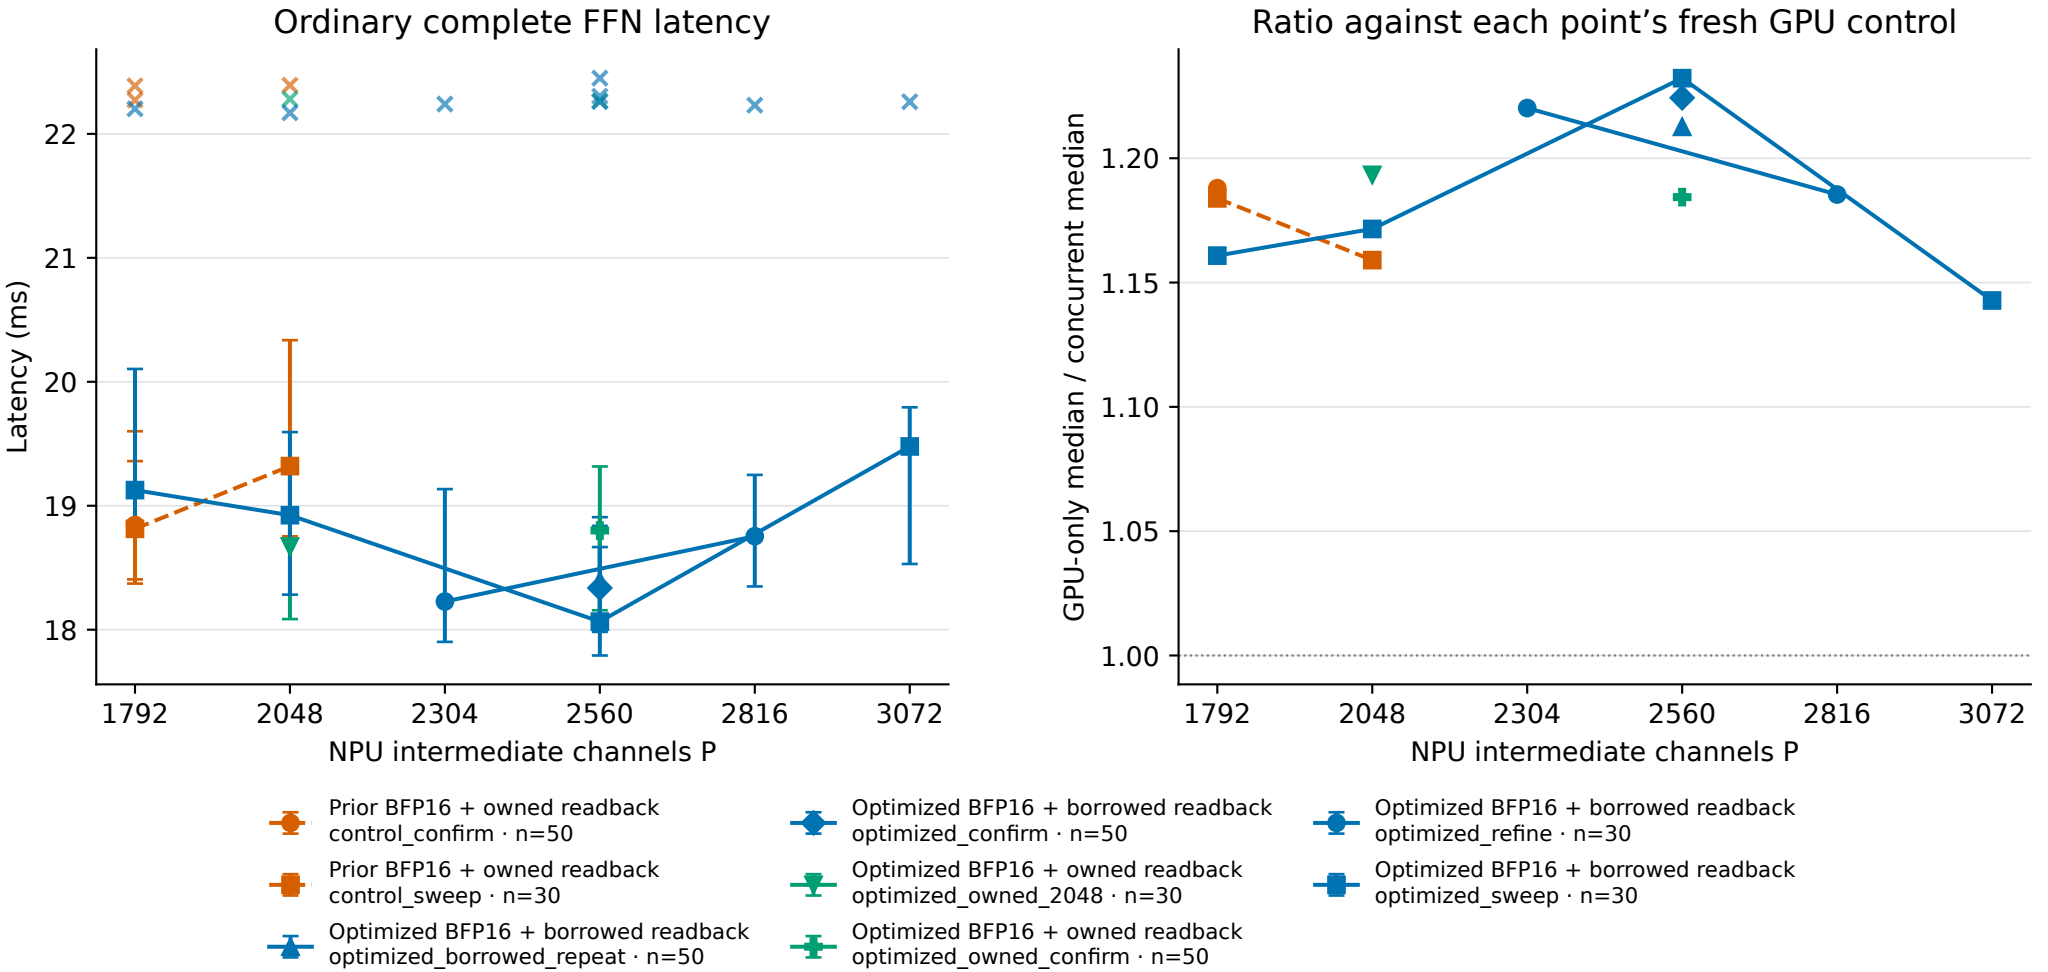

Error bars: sample p10–p90; × markers: matched GPU-only medians. Each repeated run remains separate.
Synthetic isolated FFN using HeteroMosaic GPU kernels, not native heterogeneous or full-model prefill. NPU power is only established by recorded snapshots.